

NEXUSFLOW INC.

POST-INCIDENT REVIEW (PIR) & ROOT CAUSE ANALYSIS (RCA) REPORTS

Compilation of 30 Fictional Incident Reports

CONFIDENTIAL: For Internal Use Only

Generated: May 23, 2024

INTRODUCTION

This document contains a compilation of 30 fictional Root Cause Analysis (RCA) reports for NexusFlow Inc., a global streaming and data analytics company. These reports cover a variety of scenarios across different platforms (AthenaStream, DataNexus, NexusCommerce, Internal Tools) and AWS services to serve as a learning and training tool for the Cloud Engineering team.

All incidents are fictional and created for educational purposes.

=====

INCIDENT #: NFI-2023-0010

DATE: October 26, 2023

SEVERITY: SEV-1

PLATFORM: AthenaStream

SERVICES: Amazon CloudFront, EC2

TITLE: Global Video Buffering Incident

SUMMARY: Users worldwide experienced severe buffering and failed video playback for 45 minutes.

ROOT CAUSE: A deployment of a new "recommendation engine" to an EC2 Auto Scaling Group (ASG) contained a flawed memory cache configuration. The instances consumed all available memory within minutes of launch, causing them to fail health checks. The ASG continuously terminated unhealthy instances and launched new ones, which also failed. This rapid, continuous cycling of every instance in the pool caused a thundering herd problem on the backend database, crippling the entire service.

RESOLUTION: Rolled back the deployment by attaching the previous launch template to the ASG. Implemented a scaling policy to add instances more gradually.

PREVENTION:

1. Implement canary deployments with a slower traffic shift.
2. Add stricter memory-based health checks in pre-production environments.
3. Define and test rollback procedures under load.

INCIDENT #: NFI-2023-0011

DATE: November 5, 2023

SEVERITY: SEV-2

PLATFORM: DataNexus

SERVICES: AWS Glue, S3

TITLE: ETL Pipeline Failure

SUMMARY: Nightly ETL jobs failed, causing dashboards to display stale data.

ROOT CAUSE: A Glue Job's IAM execution role did not have the `s3:ListBucket` permission on a newly created target S3 bucket. The job could write data but failed during its final step to list and verify the output, causing it to error out.

RESOLUTION: Added the missing `s3:ListBucket` permission to the IAM role. Manually triggered the failed jobs.

PREVENTION:

1. Implement a standardized IAM policy template for Glue jobs.
2. Add a pre-flight check in the CI/CD pipeline to validate required permissions against the target environments.

INCIDENT #: NFI-2023-0012

DATE: November 12, 2023

SEVERITY: SEV-2

PLATFORM: NexusCommerce

SERVICES: Amazon RDS (Aurora PostgreSQL), Lambda

TITLE: Checkout Latency

SUMMARY: Checkout process was timing out for 20% of users during a flash sale event.

ROOT CAUSE: The `orders` table in the Aurora PostgreSQL database lacked an index on the `user_id` and `created_at` columns. A critical query executed by a Lambda function during checkout performed a full table scan, which became exponentially slower as the table grew during the sale event.

RESOLUTION: Added a composite index on `(user_id, created_at)` to optimize the query.

PREVENTION:

1. Implement a process for regular database performance review and index analysis.
2. Use AWS Database Advisor and enable Performance Insights on all RDS instances.

INCIDENT #: NFI-2023-0013

DATE: November 18, 2023

SEVERITY: SEV-1

PLATFORM: Internal Tools

SERVICES: Amazon VPC, NACLs

TITLE: CRM Application Outage

SUMMARY: The entire CRM application became unreachable.

ROOT CAUSE: A network engineer attempted to block a specific malicious IP address by modifying the Network ACL (NACL). The new rule to deny the IP was placed *after* the default "allow all" rule. In the process, a typo was made in the rule number for the standard "Allow All" rule, inadvertently deleting it and replacing it with an effective "Deny All" rule.

RESOLUTION: Identified the erroneous NACL change and restored the correct allow rules.

PREVENTION:

1. Implement infrastructure as code (IaC) for all network changes (e.g., via Terraform).
2. Require a peer review for all manual NACL and Security Group changes.

INCIDENT #: NFI-2023-0014

DATE: November 22, 2023

SEVERITY: SEV-2

PLATFORM: AthenaStream

SERVICES: Amazon CloudFront

TITLE: CDN Cache Invalidation Storm

SUMMARY: Origin server load increased by 300%, causing latency for users.

ROOT CAUSE: A script designed to invalidate cached assets for a new movie release used a wildcard (`/*`) instead of a specific path prefix (e.g., `/assets/movie_123/*`). This invalidated the entire CloudFront cache for the distribution, forcing all subsequent user requests to hit the origin servers simultaneously.

RESOLUTION: The issue resolved itself as the cache gradually warmed back up. Monitored origin capacity to ensure it handled the load.

PREVENTION:

1. Implement approval workflows in CI/CD for invalidation requests containing `/*`.
2. Create a standardized script for teams to use that constructs precise invalidation paths.

INCIDENT #: NFI-2023-0015

DATE: December 1, 2023

SEVERITY: SEV-3

PLATFORM: DataNexus

SERVICES: AWS Cost Explorer, S3

TITLE: Cost Overrun Alert

SUMMARY: Unexpected \$5,000 spike in S3 costs.

ROOT CAUSE: A developer's script, which was intended to run once, was incorrectly configured and began running in an infinite loop. Each iteration uploaded several MBs of test data to an S3 bucket. Over a weekend, this resulted in millions of PUT requests and several TB of storage.

RESOLUTION: Identified and terminated the script. Implemented S3 Lifecycle Policies to automatically delete objects in the target bucket after 1 day.

PREVENTION:

1. Implement S3 Cost Allocation tags.
2. Create AWS Budgets alerts with a lower threshold.
3. Use S3 Intelligent-Tiering for unpredictable access patterns.

INCIDENT #: NFI-2023-0016

DATE: December 5, 2023

SEVERITY: SEV-2

PLATFORM: NexusCommerce

SERVICES: Amazon API Gateway, Lambda

TITLE: Failed Payment Processing

SUMMARY: Payment processing API returned 5xx errors for 15 minutes.

ROOT CAUSE: The Lambda function behind the API Gateway endpoint hit the default concurrency limit (1000). A sudden surge in traffic from a marketing campaign caused the function to throttle, rejecting additional invocation requests.

RESOLUTION: Increased the reserved concurrency limit for the critical function and configured automatic scaling.

PREVENTION:

1. Proactively set appropriate concurrency limits for functions expected to receive high traffic.
2. Use AWS Auto Scaling or configure provisioned concurrency for predictable burst traffic.

INCIDENT #: NFI-2023-0017

DATE: December 10, 2023

SEVERITY: SEV-2

PLATFORM: Internal Tools

SERVICES: AWS IAM Identity Center (AWS SSO)

TITLE: SSO Authentication Failure

SUMMARY: Employees could not log in to the AWS Management Console via SSO.

ROOT CAUSE: The permission set attached to the "ReadOnly" group was accidentally modified, removing the `sts:AssumeRole` permission. This broke the ability for users to assume the role that grants them console access.

RESOLUTION: Re-applied the correct managed policy to the permission set.

PREVENTION:

1. Define permission sets as code.
 2. Restrict modify permissions on production permission sets to a senior cloud admin group.
-

INCIDENT #: NFI-2023-0018

DATE: December 15, 2023

SEVERITY: SEV-1

PLATFORM: AthenaStream

SERVICES: AWS Elemental MediaLive

TITLE: Live Stream Failure

SUMMARY: A major live sporting event stream failed to start on time.

ROOT CAUSE: The MediaLive channel was in a "IDLE" state and needed to be manually started. The runbook for live events assumed the channel would be started via an automated Lambda function, but a configuration error in the function's event source (Amazon EventBridge) prevented the execution command from being sent.

RESOLUTION: Manually started the channel, causing a 7-minute delay to the live stream.

PREVENTION:

1. Implement a dual-check process where an engineer must confirm the channel status 1 hour before a major event.
 2. Add a CloudWatch alarm for channels that are in "IDLE" state 30 minutes before a scheduled start.
-

INCIDENT #: NFI-2023-0019

DATE: December 20, 2023

SEVERITY: SEV-1

PLATFORM: DataNexus

SERVICES: Amazon DynamoDB

TITLE: Data Loss Incident

SUMMARY: One week of user analytics data was deleted from a DynamoDB table.

ROOT CAUSE: A developer ran a script with a `DeleteItem` command in a production environment, intended for a pre-production table. The script had a hard-coded table name variable pointing to the production table ARN.

RESOLUTION: Restored the table to a point-in-time backup from 4 hours prior using PITR. The 4 hours of data between the backup and the incident were lost.

PREVENTION:

1. Enforce the use of environment-specific configuration files, never hard-coded ARNs.
 2. Implement mandatory infrastructure and tooling reviews for scripts performing destructive operations.
-

INCIDENT #: NFI-2024-0001

DATE: January 8, 2024

SEVERITY: SEV-2

PLATFORM: All

SERVICES: AWS SDK, us-west-2 Region

TITLE: Regional API Latency

SUMMARY: Applications experienced intermittent timeouts when calling AWS APIs in `us-west-2`.

ROOT CAUSE: An application's retry logic was overly aggressive. During a brief period of network congestion, the application's exponential backoff was disabled, causing it to retry failed calls instantly and repeatedly. This created a retry storm that further degraded performance.

RESOLUTION: The network congestion resolved. The application's flawed retry logic was identified and fixed.

PREVENTION:

1. Mandate the use of the latest AWS SDKs, which have built-in, sane retry mechanisms.
 2. Code reviews must include logic for API call retries and backoff.
-

INCIDENT #: NFI-2024-0002

DATE: January 15, 2024

SEVERITY: SEV-2

PLATFORM: NexusCommerce

SERVICES: Amazon ElastiCache (Redis)

TITLE: Shopping Cart Emptying

SUMMARY: User shopping carts were spontaneously emptying.

ROOT CAUSE: The Redis cluster's `maxmemory-policy` was set to `volatile-lru`. A memory pressure event caused Redis to evict keys with a TTL to make space. However, many cart items were stored without a TTL and should have been persistent. The policy was incorrectly set.

RESOLUTION: Scaled the ElastiCache node type to one with more memory. Corrected the eviction policy.

PREVENTION:

1. Define CloudFormation templates for ElastiCache to ensure consistent configuration.
 2. Implement CloudWatch alarms for memory utilization.
-

INCIDENT #: NFI-2024-0003

DATE: January 22, 2024

SEVERITY: SEV-3

PLATFORM: Internal Tools

SERVICES: AWS Backup

TITLE: Backup Failure

SUMMARY: Scheduled nightly backups for an EC2 instance failed for 5 consecutive days.

ROOT CAUSE: The EC2 instance's IAM role lacked the `aws-backup` tag. The AWS Backup resource selection was configured to automatically include only instances with that specific tag. The instance was never included in the backup plan.

RESOLUTION: Added the required tag to the instance IAM role. Manually initiated a backup to verify success.

PREVENTION:

1. Create a proactive compliance check using AWS Config to flag EC2 instances missing the backup tag.
 2. Perform periodic recovery drills.
-

INCIDENT #: NFI-2024-0004

DATE: January 30, 2024

SEVERITY: SEV-1

PLATFORM: AthenaStream

SERVICES: Amazon Route 53

TITLE: DNS Resolution Failure

SUMMARY: The primary domain `athenastream.com` became unresolvable.

ROOT CAUSE: During a DNS record update, a misconfigured CI/CD pipeline accidentally deleted the entire hosted zone's NS (Name Server) records instead of just updating the A-record alias.

RESOLUTION: Recovered the NS records from a Terraform state file and reapplied them to the hosted zone. Full DNS propagation took ~30 minutes.

PREVENTION:

1. Implement manual approval gates in the deployment pipeline for changes to critical DNS records (especially NS and SOA).
 2. Use Terraform `prevent_destroy` lifecycle flags on the hosted zone resource.
-

INCIDENT #: NFI-2024-0005

DATE: February 5, 2024

SEVERITY: SEV-2

PLATFORM: DataNexus

SERVICES: Amazon Redshift

TITLE: Redshift Query Performance Degradation

SUMMARY: Business intelligence queries were running 10x slower than usual.

ROOT CAUSE: A large, unoptimized `COPY` operation from S3 loaded a massive dataset without sorting it on the intended sort key. This led to massive zone maps being ineffective and required Redshift to scan almost all blocks for subsequent queries.

RESOLUTION: Ran a `VACUUM FULL` and `ANALYZE` on the affected table to re-sort the rows and update statistics.

PREVENTION:

1. Implement data loading best practices: use sort keys, distribute keys, and break large copies into smaller files.
 2. Schedule regular maintenance operations during off-peak hours.
-

INCIDENT #: NFI-2024-0006

DATE: February 14, 2024

SEVERITY: SEV-2

PLATFORM: DataNexus

SERVICES: Amazon Kinesis Data Firehose

TITLE: Kinesis Firehose Delivery Stream Backpressure

SUMMARY: Real-time clickstream data was delayed by over 2 hours.

ROOT CAUSE: The Kinesis Firehose delivery stream was configured to transform records using a Lambda function. A code update to the Lambda function introduced an error that

caused it to timeout on 30% of records. This throttling created backpressure, drastically reducing the stream's throughput.

RESOLUTION: Identified the failing Lambda function from CloudWatch Logs. Rolled back the Lambda function code to the previous stable version. The Firehose stream cleared its backlog automatically.

PREVENTION:

1. Implement more robust error handling and logging in data transformation Lambdas.
2. Test Lambda function changes against a sample of production data in a staging Firehose stream.

INCIDENT #: NFI-2024-0007

DATE: February 21, 2024

SEVERITY: SEV-1

PLATFORM: AthenaStream

SERVICES: Amazon Route 53 (Failover Routing Policy)

TITLE: Regional Failover Test Failure

SUMMARY: A scheduled DR drill failed; traffic did not fail over to the secondary region.

ROOT CAUSE: The health check for the primary region's endpoint was configured with an overly high threshold (5 consecutive failures). The simulated outage was not long enough to trigger the health check to fail.

RESOLUTION: Adjusted the health check to fail after 2 consecutive failures of 30 seconds each.

PREVENTION:

1. Document and test all health check configurations as part of the DR runbook.
2. Use a blue/green deployment model for critical health check changes.

INCIDENT #: NFI-2024-0008

DATE: February 28, 2024

SEVERITY: SEV-2

PLATFORM: NexusCommerce

SERVICES: AWS WAF

TITLE: WAF False Positives Blocking Checkout

SUMMARY: Legitimate users in specific geographic regions were blocked during checkout.

ROOT CAUSE: A new WAF rule was deployed to block traffic from a country code associated with a high volume of fraud. The rule was too broad and did not include exceptions for known-good IP ranges.

RESOLUTION: Modified the WAF rule to add allow-list exceptions for trusted ASNs and IP ranges.

PREVENTION:

1. Deploy new WAF rules in "Count" mode first to analyze potential false positives.
2. Create a formal review process for geo-blocking rules.

INCIDENT #: NFI-2024-0009

DATE: March 7, 2024

SEVERITY: SEV-2

PLATFORM: Internal Tools

SERVICES: AWS Secrets Manager

TITLE: Secrets Rotation Failure

SUMMARY: A critical internal application failed to start after an automated secrets rotation.

ROOT CAUSE: The application was configured to cache the database credential from Secrets Manager in its environment variables upon startup. It did not have logic to re-fetch the secret from Secrets Manager if it became invalid.

RESOLUTION: Restarted the application, forcing it to fetch the new secret. Updated the application code to use the Secrets Manager SDK to fetch the secret dynamically on every use.

PREVENTION:

1. Audit all applications using Secrets Manager to ensure they handle rotation correctly.
 2. Use the built-in Secrets Manager JDBC/ODBC drivers for database connections where possible.
-

INCIDENT #: NFI-2024-0010

DATE: March 14, 2024

SEVERITY: SEV-3

PLATFORM: DataNexus

SERVICES: Amazon S3 (Replication)

TITLE: S3 Cross-Region Replication Delay

SUMMARY: Data scientists reported a 12-hour delay in data availability in the secondary region for analytics.

ROOT CAUSE: S3 Replication was configured for a bucket with existing data. A new lifecycle policy was added to archive objects to S3 Glacier Instant Retrieval. Replication does not process objects that have been transitioned to a storage class other than Standard, Standard-IA, or One Zone-IA after the replication configuration was made.

RESOLUTION: Manually synced the affected objects using AWS CLI `sync` command. Modified the process: data is now replicated immediately after upload, and the lifecycle policy is applied only in the destination region.

PREVENTION:

1. Thoroughly review S3 replication and lifecycle policy interactions before deployment.
 2. Monitor the S3 replication metrics (SourceBucketOps and ReplicationLatency) in CloudWatch.
-

INCIDENT #: NFI-2024-0011

DATE: March 21, 2024

SEVERITY: SEV-2

PLATFORM: AthenaStream

SERVICES: Amazon Elasticsearch Service (OpenSearch Service)

TITLE: Elasticsearch Cluster Red Status

SUMMARY: Search functionality on the platform was slow and returning partial results.

ROOT CAUSE: One of the three data nodes in the OpenSearch cluster failed due to an underlying hardware issue. The cluster health turned "red" because one primary shard and its replicas were unassigned.

RESOLUTION: AWS service health automatically replaced the failed node. The cluster recovered to a "green" state as shards were reallocated.

PREVENTION:

1. Implement CloudWatch alarms for cluster status ("red" and "yellow") and node count.
 2. Consider using Multi-AZ deployment for critical clusters for higher availability.
-

INCIDENT #: NFI-2024-0012

DATE: March 28, 2024

SEVERITY: SEV-2

PLATFORM: NexusCommerce

SERVICES: AWS Step Functions

TITLE: Step Functions Execution Limit Hit

SUMMARY: The order fulfillment workflow stopped processing new orders.

ROOT CAUSE: The Step Functions state machine had a concurrency limit of 1000 simultaneous executions. A surge in orders, combined with a downstream vendor API being slow, caused the limit to be reached.

RESOLUTION:

1. Increased the execution limit as a short-term fix.
2. Worked with the vendor to resolve their API latency.

PREVENTION:

1. Monitor Step Functions cloudwatch metrics for `ExecutionThrottled`.
2. Architect for asynchronous processing or implement a queueing mechanism in front of the state machine.

INCIDENT #: NFI-2024-0013

DATE: April 4, 2024

SEVERITY: SEV-3

PLATFORM: Internal Tools

SERVICES: AWS CloudFormation

TITLE: CloudFormation Stack Rollback Failure

SUMMARY: A CloudFormation stack update failed and could not roll back, entering the `UPDATE_ROLLBACK_FAILED` state.

ROOT CAUSE: The update added a new EC2 instance with an associated EIP. The rollback process tried to delete this EIP but failed because it was associated with the EC2 instance.

RESOLUTION: Manually disassociated the EIP, then successfully executed a "Continue Update Rollback" in the CloudFormation console.

PREVENTION:

1. Use CloudFormation deletion policies more carefully.
2. For complex resources, consider custom resources or breaking stacks into smaller, more manageable pieces.

INCIDENT #: NFI-2024-0014

DATE: April 11, 2024

SEVERITY: SEV-2

PLATFORM: DataNexus

SERVICES: Amazon EKS

TITLE: EKS Control Plane Connectivity Issues

SUMMARY: Engineers could not use `kubectl` to list pods or deploy new services for a 20-minute period.

ROOT CAUSE: The EKS cluster endpoint was configured as public but restricted to the corporate IP range. The corporate NAT gateway's public IP address changed unexpectedly, cutting off all access to the EKS API server.

RESOLUTION: Updated the EKS cluster's security configuration to allow the new NAT gateway IP address.

PREVENTION:

1. Configure a VPC-native EKS cluster with a private endpoint and use a bastion host or VPN for access.
2. Implement a Lambda function to monitor and automatically update security groups on IP change events.

INCIDENT #: NFI-2024-0015

DATE: April 18, 2024

SEVERITY: SEV-2

PLATFORM: AthenaStream

SERVICES: EC2 Instance Metadata Service (IMDSv1)

TITLE: High EC2 Metadata API Calls

SUMMARY: Some application instances experienced network latency.

ROOT CAUSE: An application was incorrectly making a call to the EC2 metadata service (IMDSv1) *on every web request* to retrieve the instance ID for logging purposes.

RESOLUTION: Refactored the application to retrieve the instance ID once at startup and cache it, instead of querying on every request.

PREVENTION:

1. Code reviews should flag inappropriate use of instance metadata.
2. Consider enforcing IMDSv2, which has a more secure hop limit and request flow.

INCIDENT #: NFI-2024-0016

DATE: April 25, 2024

SEVERITY: SEV-2

PLATFORM: NexusCommerce

SERVICES: Amazon RDS (Read Replica)

TITLE: RDS Read Replica Lag

SUMMARY: Users reported seeing stale data on their account pages, which reads from a read replica.

ROOT CAUSE: A long-running analytical query was executed directly on the read replica for reporting. The query locked tables and consumed excessive IOPS, causing significant replication lag (over 5 minutes).

RESOLUTION: Killed the offending query on the replica. Established a dedicated reporting replica separate from the one used for low-latency application reads.

PREVENTION:

1. Implement user permission boundaries to prevent analysts from running ad-hoc queries on production replicas.
2. Use Amazon Aurora to leverage dedicated reader endpoints that can better isolate workload types.

INCIDENT #: NFI-2024-0017

DATE: May 2, 2024

SEVERITY: SEV-3

PLATFORM: Internal Tools

SERVICES: AWS Config

TITLE: Config Rule Non-Compliance

SUMMARY: A compliance audit revealed unencrypted S3 buckets that were supposed to be enforced by AWS Config.

ROOT CAUSE: A custom AWS Config rule designed to check for S3 encryption was written to only apply to buckets tagged with `env=prod`. Several new production buckets were created without this specific tag.

RESOLUTION: Applied the correct tag to the buckets. Updated the Config rule logic to also identify production buckets by a naming convention prefix.

PREVENTION:

1. Use Service Control Policies (SCPs) to enforce mandatory tags at the AWS Organization level.
 2. Design Config rules to be based on resource types, not tags, where possible.
-

INCIDENT #: NFI-2024-0018

DATE: May 9, 2024

SEVERITY: SEV-2

PLATFORM: DataNexus

SERVICES: AWS Lambda

TITLE: Lambda Function Versioning Conflict

SUMMARY: A data processing Lambda function began emitting data in an incorrect format, breaking downstream systems.

ROOT CAUSE: An alias (`LIVE`) was pointing to the `\$LATEST` version of the Lambda function instead of a specific, stable numbered version. A developer accidentally deployed unfinished code directly to `\$LATEST`.

RESOLUTION: Re-pointed the `LIVE` alias back to the previous stable version to immediately revert the change.

PREVENTION:

1. Never point production aliases to `\$LATEST`.
 2. Implement a deployment process that publishes a new version and updates the alias atomically upon successful testing.
-

INCIDENT #: NFI-2024-0019

DATE: May 16, 2024

SEVERITY: SEV-3

PLATFORM: AthenaStream

SERVICES: Amazon CloudWatch Logs

TITLE: CloudWatch Logs Insights Query Timeout

SUMMARY: Engineers were unable to run Logs Insights queries to debug a production issue.

ROOT CAUSE: A misconfigured application logger was writing the entire HTTP request/response payload (including large base64-encoded images) to CloudWatch Logs for every API call.

RESOLUTION: Fixed the application logging configuration to only log metadata, not payloads. Archived the existing large log streams to S3.

PREVENTION:

1. Implement log-level standards and reviews.
 2. Use structured logging (JSON) and filter patterns to avoid logging unnecessary data.
-

INCIDENT #: NFI-2024-0020

DATE: May 23, 2024

SEVERITY: SEV-3

PLATFORM: All

SERVICES: AWS Budgets, SNS, Email

TITLE: Billing Alarm Notification Failure

SUMMARY: The monthly billing alarm did not notify the finance team, leading to a late surprise about costs.

ROOT CAUSE: The SNS topic for billing alerts was configured to send email to a distribution list that was decommissioned without updating the SNS subscription, causing all notification emails to bounce.

RESOLUTION: Updated the SNS topic subscription to a new, valid email address. Verified the alarm notification worked.

PREVENTION:

1. Create a runbook for decommissioning resources that includes checking for AWS dependencies like SNS subscriptions.
2. Use a shared Slack channel integrated with SNS for critical alerts instead of email alone.

=====

=====

****END OF DOCUMENT****